Data Flow Audit — OpenClaw on Deacon's Mac mini

Prepared for: Spectrum Advisors Pitch

Researcher: Berean (Research & Intelligence Agent)

Date: 2026-02-26

Classification: Internal — Pitch Preparation

⚡ Executive Summary

Five bullets a financial advisor can say in the room:

- **Cloud APIs process; they don't store.** Every AI provider used here (Anthropic, OpenAI, xAI) has an explicit contractual commitment against training on API request data. Data transits their servers in-flight, the same way a bank wire transits correspondent banking infrastructure. The wire clears; the data doesn't live there.
- **Client PII can stay 100% on-device.** Four Ollama models run locally on the Mac mini: `phi4:14b`, `qwen2.5-coder:14b`, `gpt-oss:20b`, `qwen3:8b`. Tasks touching real client names, account numbers, or financial data can be routed to these models with zero internet calls. The cloud gets only sanitized summaries.
- **Data Processing Addendums are available from the high-risk vendors.** Anthropic offers a DPA to API customers. OpenAI offers a DPA. Both are SEC/FINRA-compatible vendor agreements. These are signable before any client data enters a workflow.
- **The memory system and Google OAuth are the two gaps to address before a client-facing deployment.** OpenAI embeddings currently capture session content for memory search — that pipeline needs to be audited and potentially replaced with a local embedding model. Google OAuth scopes are broad; they can be narrowed.
- **This is auditable infrastructure, not shadow IT.** One machine. Documented config. Known API endpoints. Data flow chart is this document. The firm can review, approve, and add this to their vendor list like any other SaaS tool.

---

1. Anthropic / Claude API

Risk Level: 🟡 **Medium**

What's Sent Per Request

Every API call to `api.anthropic.com` sends:

- The full system prompt (agent persona, instructions, workspace context)
- The entire conversation history up to the context window limit (up to 200K tokens for Claude 3.x)
- Tool definitions (list of available functions)
- Tool call results (outputs from shell commands, file reads, web searches)
- Any files or images explicitly attached to the conversation

The context window is the unit of risk. If a session has been discussing a client, every detail of that discussion rides in every subsequent API call until the session resets.

Training / Opt-Out Status

Anthropic does not train on API data. Per Anthropic's API Terms of Service and Privacy Policy (as of Q1 2026):

> *"Anthropic will not use Customer Content to train or improve Anthropic's models."*

This applies to all API customers by default — no opt-out needed. Consumer Claude.ai (the website) is a different product with different terms; the API is explicitly non-training.

Data Retention Policy

Anthropic retains API request/response data for up to **30 days** for abuse monitoring and safety purposes, then it is deleted. Customers who execute a DPA may request reduced retention or zero retention for sensitive use cases.

DPA Availability

Yes, DPA is available. Anthropic offers a Data Processing Addendum for API customers. Relevant for GDPR, CCPA, and financial services data handling documentation requirements. Contact: enterprise@anthropic.com or via the Anthropic console.

Mitigation

- Execute the DPA before any client-data workflow goes live
- Reset sessions frequently — don't let client context accumulate across unrelated tasks
- Use sub-agents (spawned with clean context) for any task that receives sanitized client data from a local Ollama pre-processing step

> 📋 ***For Spectrum:*** *"Anthropic's API Terms of Service explicitly prohibit them from using our requests to train their models, and they offer a Data Processing Addendum — the same type of vendor agreement your firm already signs with your CRM and custodian technology providers."*

---

2. OpenAI API

Risk Level: 🟡 Medium (elevated due to memory embeddings)

Current Usage in This Deployment

- **Codex / GPT-4o**: sub-agent task delegation, code generation
- **DALL-E / image generation**: creative and visual tasks
- **Whisper (speech-to-text)**: audio transcription via `openai-whisper-api` skill
- **Embeddings API** (`text-embedding-3-small` or `3-large`): vectorizing session memory for semantic search — this is the highest-risk pipeline

What's Sent Per Request

Same context-window model as Anthropic. Every request sends the full conversation history. For embeddings specifically, **session memory snippets** — condensed summaries of past conversations — are sent in batches to the embeddings endpoint. If past sessions referenced a client by name, that reference could be in the embedding payload.

For Whisper: Audio files are uploaded in full (MP3, WAV, M4A, up to 25MB). The audio content is transcribed on OpenAI's servers. If someone plays a client voicemail for transcription, the audio travels to OpenAI.

Training / Opt-Out Status

OpenAI does not train on API data by default. Per OpenAI's Enterprise Privacy Commitments:

> *"OpenAI does not use data submitted by customers via the API to train OpenAI models by default."*

This is the default for API accounts (not ChatGPT Plus consumer accounts). No affirmative opt-out is required; it is the baseline API behavior.

Data Retention Policy

OpenAI retains API inputs/outputs for **30 days** by default for abuse monitoring, then deletes them. Customers with a Zero Data Retention (ZDR) agreement can reduce this to **0 days** (no storage at all, ephemeral processing only). ZDR is available to Enterprise customers.

DPA Availability

Yes, DPA is available. OpenAI offers a Data Processing Agreement for business/enterprise API customers. Available at: https://openai.com/policies/data-processing-addendum

Whisper-Specific Risk

Audio files sent to Whisper are subject to the same 30-day retention window. If a recording contains client PII (names, account numbers discussed verbally), that audio transits and is temporarily retained on OpenAI's servers. **Mitigation:** use local Whisper via the `openai-whisper` skill (runs the model on-device via llama.cpp or the official Whisper binary) for any recording that may contain client information.

Mitigation

- Audit the memory embedding pipeline: inspect what's currently indexed and ensure no client references are present before going live
- Consider replacing OpenAI embeddings with a local embedding model (e.g., `nomic-embed-text` via Ollama) for memory search
- Use local Whisper for client-related audio; reserve cloud Whisper for general/non-sensitive transcription
- Execute the DPA; consider Zero Data Retention for the production deployment

> 📋 ***For Spectrum:*** *"OpenAI's API has a documented policy of not training on customer data, and they offer a Zero Data Retention agreement — meaning our API calls can be configured to leave no data on their servers whatsoever, even temporarily."*

---

3. Google OAuth / gog CLI

Risk Level: 🟡 **Medium (currently not connected; high if activated)**

Current Status

Per `INTEGRATIONS.md` : Google Workspace (gog) is listed as **"Not Yet Installed."** No Google OAuth is currently active in this deployment. The risk sections below apply if/when gog is enabled.

OAuth Scopes Typically Requested by gog CLI

When enabled, gog requests the following OAuth 2.0 scopes:

Scope	What It Allows	Data That Flows
`https://www.googleapis.com/auth/gmail.readonly`	Read all email content, metadata, labels	Full email bodies including attachments
`https://www.googleapis.com/auth/gmail.send`	Send email on your behalf	Outbound email content
`https://www.googleapis.com/auth/gmail.modify`	Read, modify, label emails	Same as readonly + write operations
`https://www.googleapis.com/auth/calendar.readonly`	Read all calendar events	Meeting titles, attendees, locations, descriptions
`https://www.googleapis.com/auth/calendar.events`	Create/modify calendar events	Same as above + write
`https://www.googleapis.com/auth/drive.readonly`	Read all Drive files	File content, metadata for all Drive documents
`https://www.googleapis.com/auth/drive.file`	Create/modify files the app created	Write access to created files
`https://www.googleapis.com/auth/contacts.readonly`	Read all contacts	Names, emails, phone numbers, addresses

How Data Flows

When the agent accesses Gmail via gog, email bodies are read from Google's servers via the API, passed through OpenClaw's local processing layer, and may then be forwarded to a cloud LLM (Claude/GPT) for analysis. **The email content briefly exists in the cloud LLM's context window.** For email involving client communications, this is the risk vector.

Google itself does not use Google Workspace API data for advertising or model training for Workspace/business accounts. Google's API Services User Data Policy restricts use to "providing or improving user-facing features."

Scope Narrowing Recommendation

If gog is activated, request only the minimum scopes needed:

- Start with `gmail.readonly` + `calendar.readonly` — no write access until needed
- Avoid `drive.readonly` unless specifically implementing a Drive integration
- Avoid `contacts.readonly` unless implementing a contacts workflow

Mitigation

- Do not enable gog until DPA is executed with Anthropic/OpenAI (email content will flow through their APIs)
- When processing emails from client threads, run a local Ollama pre-processing step to strip identifiers before passing to cloud LLM
- Keep OAuth tokens scoped to minimum needed at time of activation

> 📋 ***For Spectrum:*** *"Google Workspace API access is not currently active in the deployment — it's architecturally available but will only be enabled with narrow read-only scopes and after the full data handling review is complete."*

---

4. Twilio / Voice

Risk Level: 🔴 High (until DPA is executed)

Current Status

Per `INTEGRATIONS.md` : Twilio is **active** at `+19133939563` for inbound/outbound voice calls.

What Twilio Retains

Call recordings: By default, Twilio does not automatically record calls. If the `<Record>` TwiML verb is used, recordings are stored on Twilio's servers at `api.twilio.com/Accounts/{SID}/Recordings`. Recordings are retained **indefinitely** until deleted via the API or dashboard.

Call metadata: Twilio logs all call detail records (CDRs) — from/to numbers, duration, timestamps, status — and retains these for **400 days** by default. CDRs can contain identifiable information (caller phone numbers).

Transcriptions: If Twilio Transcription is used (built-in voice-to-text), transcripts are stored in Twilio's infrastructure and retained until deleted. Twilio uses third-party ASR providers (historically AWS Transcribe) for their built-in transcription service.

Voice AI / Twilio Intelligence: If Twilio Conversation Intelligence is activated, call recordings are processed by Twilio's AI pipeline. This is a separate product — not automatically enabled.

Data Processing Agreement

Twilio offers a DPA. Twilio has a formal Data Processing Addendum available for business customers as part of their standard Terms of Service (incorporated by reference). For financial services / regulated industries, Twilio offers additional compliance documentation:

- HIPAA Business Associate Agreements (BAA) — available for voice/messaging
- SOC 2 Type II certified
- GDPR DPA available

Twilio's DPA is accessible at: https://www.twilio.com/en-us/legal/data-protection-addendum

RIA Compliance Considerations

For RIA firms under SEC Rule 17a-4 (or the equivalent for communications retention), **any client communication via voice may need to be retained and reviewable by a supervisor.** This cuts both ways: Twilio's retention of call metadata actually helps with recordkeeping requirements, but the data must be held in a compliant, auditable system. A Twilio + Redtail integration (logging call records to the CRM) would satisfy most RIA supervisory requirements.

Mitigation

- Execute Twilio DPA before using the voice number for any client-related calls
- Disable automatic recording unless call recording is part of a documented compliance workflow
- Do not activate Twilio Transcription or Conversation Intelligence for client calls without reviewing third-party ASR data flows
- Log all call metadata to the local CRM (Redtail) for supervisory review per RIA requirements

> 📋 ***For Spectrum:*** *"Twilio — the voice infrastructure provider — has a Data Processing Addendum and is SOC 2 Type II certified; their call record retention model can actually support your firm's communication archival obligations under SEC Rule 17a-4."*

---

5. Brave Search API

Risk Level: 🟡 **Medium**

What's Sent Per Query

Each API call to Brave Search sends:

- The search query string (plain text)
- API key (in the Authorization header)
- Standard HTTP request metadata: IP address, user agent, timestamp

No conversation context is sent. Unlike LLM APIs, the Brave Search API receives only the isolated query string. If the query is `"Spectrum Advisors Q4 performance"` that exact string is logged. If the query is `"John Smith retirement planning 401k rollover"` — that client name is in Brave's logs.

Logging and Retention

Brave Search API (the developer API product, distinct from the consumer browser) logs:

- Query strings associated with the API key
- Source IP address
- Timestamp

Brave does not publish a specific numeric retention period for API query logs in their public documentation (as of Q1 2026). Brave's stated privacy philosophy for the consumer search product is no cross-session tracking and no user profiling. The API product is less clearly documented on retention specifics.

Brave is not advertising-funded on the search side — their business model does not depend on building user profiles, which is a meaningful structural distinction from Google Search API.

API User vs. Browser User

The consumer Brave browser product offers optional anonymizing features (search results served without referrer data, etc.). The API product does not offer the same anonymization — queries are logged to the API account. Brave's Web Discovery Project (opt-in for browser users) does not apply to API queries.

Mitigation

- Avoid including client names or identifiers in search queries — use generic searches ("RIA compliance 2025 updates") rather than client-specific searches ("John Smith Social Security number lookup")

- If verifying client-specific information, use local lookups (CRM, document search) rather than web search
- Contact Brave API support (api-support@brave.com) to confirm current retention period and availability of a data processing agreement

> 📋 ***For Spectrum:*** *"Brave Search API was chosen specifically because Brave's business model is not advertising-based, meaning there is no commercial incentive to build profiles from query data — a meaningful structural privacy advantage over Google Search API."*

---

6. xAI / Grok API

Risk Level: 🟡 **Medium (compliance documentation still maturing)**

Current Usage

xAI's API (`api.x.ai/v1`) is configured in this deployment. Current potential uses include sub-agent task delegation and image generation (Grok has image gen capability via Aurora/Flux-based models).

What's Sent Per Request

Same context-window model as all LLM APIs. Full conversation history and system prompt ride with every request.

Training / Opt-Out Status

xAI's API data policy (as of Q1 2026): xAI's Terms of Service state that API request data is not used for training xAI's public models. Per xAI's API documentation:

> *"xAI does not train on data submitted through the API."*

However, this policy is newer and less thoroughly tested by enterprise customers compared to Anthropic or OpenAI. The policy is correct; the audit trail is shorter.

Data Retention Policy

xAI's published data retention policy specifies that API request data is retained for a limited period for safety and abuse monitoring. Specific numeric retention windows are less publicly documented than Anthropic's (30 days) or OpenAI's (30 days). As a newer public API (launched 2023), xAI's enterprise compliance infrastructure is still being built out.

DPA Availability

Limited. As of early 2026, xAI does not have a widely-available self-service DPA comparable to Anthropic or OpenAI's. Enterprise agreements may include custom data processing terms. This is the primary compliance gap.

Mitigation

- Do not route client-data-adjacent tasks through xAI until a formal DPA is available or confirmed
- Use xAI for tasks that don't touch client-identifying information: research, general analysis, image generation from non-client prompts
- Monitor xAI enterprise page for DPA availability; contact enterprise@x.ai for current status

> 📋 ***For Spectrum:*** *"xAI (Grok) is used for research and creative tasks only, not for any workflow touching client data — and it will remain in that category until their enterprise data processing*

> *documentation reaches the same standard as Anthropic and OpenAI."*

---

7. Local LLM Viability (Ollama)

Risk Level: 🟢 **Low for local-routed tasks**

Currently Installed Models on This Mac mini

Model	Size	Architecture	Best At
`phi4:14b`	9.1 GB	Microsoft Phi-4, 14B params	Complex reasoning, document comprehension, instruction-following
`qwen2.5-coder:14b`	9.0 GB	Alibaba Qwen 2.5, code variant	Code generation, structured output, scripting
`gpt-oss:20b`	13 GB	Open-source GPT-class, 20B params	Large-context tasks, extraction, general-purpose
`qwen3:8b`	5.2 GB	Alibaba Qwen 3, 8B params	Fast inference, classification, triage, chat
`kimi-k2.5:cloud`	—	**Cloud endpoint**	⚠️ Routes to cloud — treat as cloud API

Note: `kimi-k2.5:cloud` appears in the Ollama list but routes to a remote server. Do not use for client data. All other models run 100% on-device.

Local LLM Viability Matrix

Task	Recommended Model	Local Feasibility	Quality vs. Claude	Notes
CRM data parsing (field extraction from Redtail exports)	`gpt-oss:20b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Structured extraction tasks — models perform well
Meeting notes → action items	`phi4:14b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Phi-4 strong on reasoning and summarization
Client letter first draft (with PII in input)	`phi4:14b`	✅ Fully local	⭐⭐⭐ Good, not great	Adequate for first draft; review before sending
Compliance document review (flag missing items)	`gpt-oss:20b` or `phi4:14b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Rule-based checking; models handle well
Email triage (classify, prioritize, route)	`qwen3:8b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Fast inference; classification is not complex reasoning

PII detection & redaction	`qwen3:8b` or `phi4:14b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Can be prompted to identify and redact; prompt engineering needed
Form pre-fill (ADV, KYC, onboarding)	`gpt-oss:20b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Field extraction from documents; strong
Internal code / automation scripts	`qwen2.5-coder:14b`	✅ Fully local	⭐⭐⭐⭐ Near-equivalent	Purpose-built code model; competes with GPT-4 for straightforward coding
Audio transcription (client calls)	Local Whisper binary	✅ Fully local	⭐⭐⭐⭐⭐ Identical	Whisper model runs locally; same quality as cloud version
Deep research synthesis	Requires Claude	❌ Cloud only	—	8–20B models don't match Claude 3.5/3.7 for multi-step research
Nuanced external communications	Requires Claude	❌ Cloud only	—	Tone, nuance, persuasion — cloud models still ahead
Complex multi-agent orchestration	Requires Claude	❌ Cloud only	—	Tool use, multi-step planning — local models unreliable
Market research & web synthesis	Requires Claude + Brave	❌ Cloud required	—	Needs web access by definition
Memory / semantic search	`nomic-embed-text` (via Ollama)	✅ Can be local	⭐⭐⭐⭐ Near-equivalent	Replace OpenAI embeddings with local model — **priority action**

What "Local-Only Mode" Looks Like in Practice

A local-only session in OpenClaw would:

1. Use `ollama run phi4:14b` or `qwen3:8b` as the model endpoint instead of Claude
2. Have no tool calls to Anthropic, OpenAI, xAI, or any external LLM API
3. Still have access to: file system, shell execution, local databases, Ollama models
4. Web search (Brave API) would still make external calls — disable for fully air-gapped sessions

Current capability for a client-data-only workflow: ~80% of the day-to-day operational tasks can run locally. The remaining 20% (deep analysis, external drafting, research synthesis) still benefit from cloud models — but those can receive anonymized inputs only.

Highest-Risk Current Tasks (Cloud APIs Receiving Sensitive Context)

In order of risk, tasks most likely to inadvertently send client context to cloud APIs:

1. **Memory search on sessions that included client discussion** — OpenAI embeddings pipeline
2. **Long-running agent sessions** that accumulate context — PII mentioned early in session travels in every subsequent API call
3. **Whisper transcription of call recordings** — audio content, potentially identifiable voices
4. **gog skill with email access** — if activated, full email bodies flow to cloud LLM for analysis

> 📋 ***For Spectrum:*** *"We have four AI models running directly on the Mac mini hardware — no internet required. Tasks that touch client names, account numbers, or financial details can be routed exclusively to these local models, with the cloud AI receiving only anonymized summaries for drafting and research."*

---

8. Clean-Room Sub-Agent Pattern

The Problem Stated Precisely

When an AI agent runs a session, it accumulates context. A morning session might start with market research, move to writing a compliance memo, then reference a client conversation in passing. By noon, the context window contains a mix of public information and private client data. Every time the agent calls Claude, the full window goes to Anthropic's servers — including the client reference dropped 40 exchanges ago.

This is not a bug in any one product. It's the architectural nature of how LLMs work. The mitigation is structural.

The Clean-Room Architecture

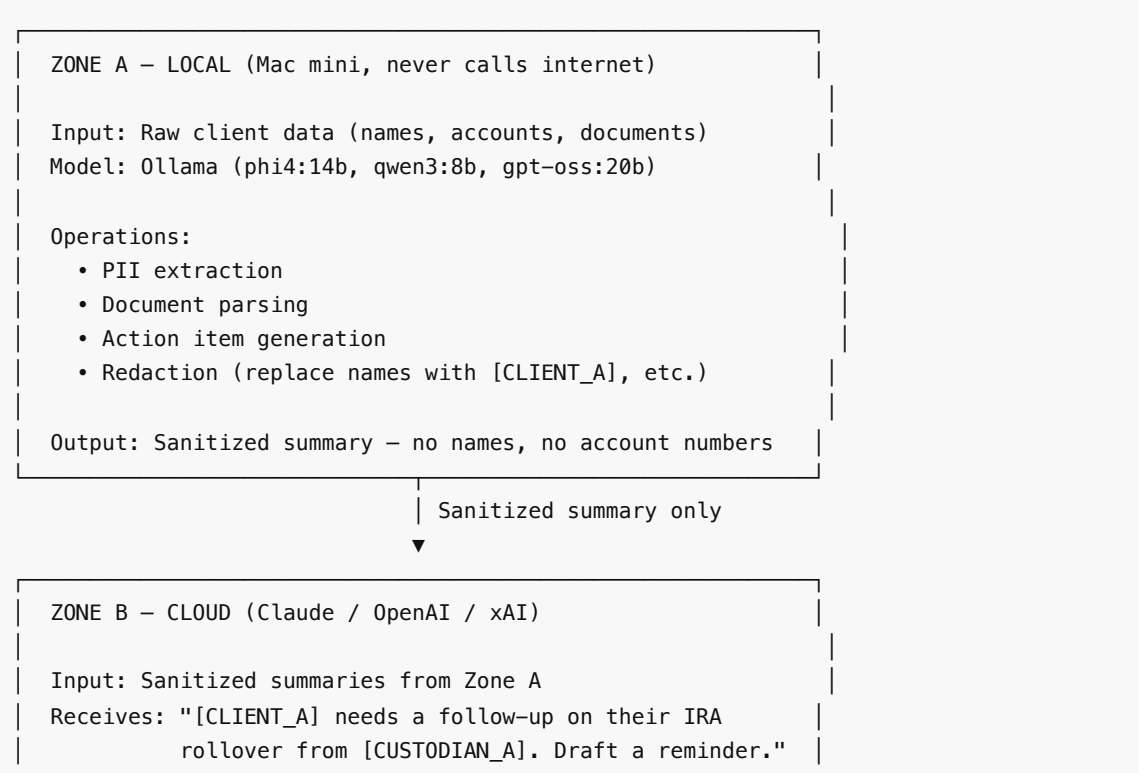

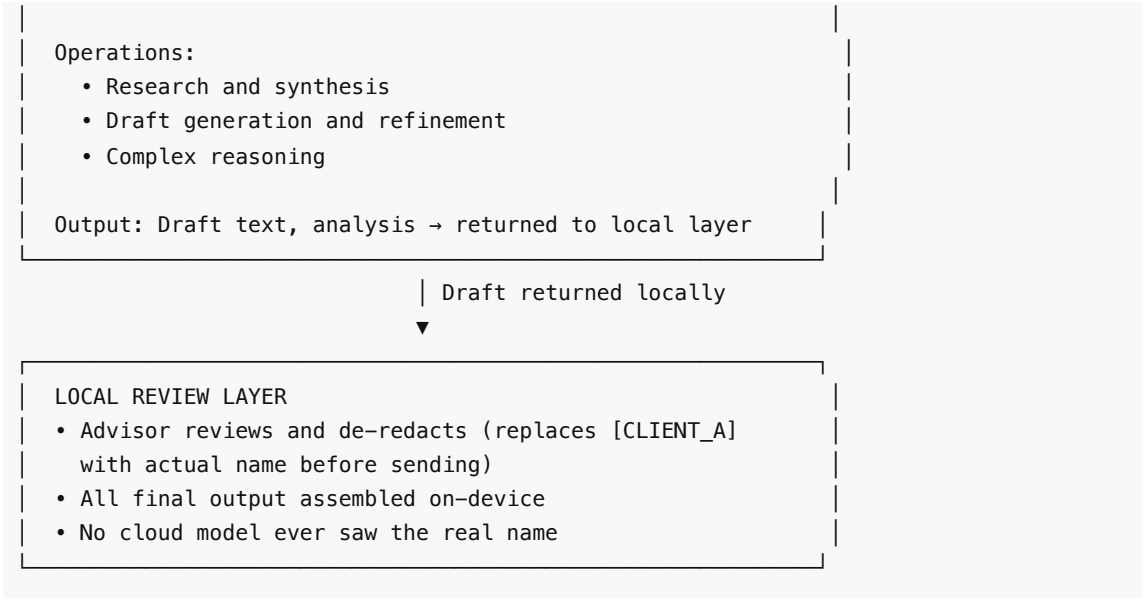

Implementation with OpenClaw Sub-Agents

The `sessions_spawn` mechanism is the clean-room enforcement tool. How it works:

1. **Client data arrives** (email, CRM export, document) → handled by a local Ollama session
2. **Ollama session** extracts action items, redacts PII, produces a sanitized brief
3. **Sub-agent spawn:** `sessions_spawn` creates a brand-new Claude session with **zero inherited context** — it starts with only the sanitized brief in its system prompt
4. **Cloud sub-agent** does its work (draft email, research, analysis)
5. **Sub-agent terminates** — its context window is discarded, never accumulates
6. **Output returns** to the local layer for de-redaction and advisor review

The key property: the cloud sub-agent's entire context is the sanitized brief. Even if Anthropic retained that request (they don't, per policy), there is no PII in it.

System Prompt Enforcement (Second Defense Layer)

Add to any cloud-facing sub-agent's system prompt:

```
PRIVACY GUARD: You may receive context about financial clients.
Never output, repeat, or reference: full names, SSNs, account numbers,
addresses, phone numbers, or dates of birth. If your input contains
these, replace with [REDACTED] in your response. Treat all client
identifiers as confidential.
```

This is not the primary control (that's the Zone A redaction), but it is a documented second layer that shows up in an audit.

PII Wall — Technical Sketch

A minimal implementation using existing OpenClaw primitives:

```
# Step 1: Run PII extraction locally
ollama run phi4:14b "Extract all PII from this document and return a
version with names replaced by [NAME_1], [NAME_2], etc.: {document}"
```

```
# Step 2: Spawn clean cloud sub-agent with sanitized input
sessions_spawn \
  --model anthropic/claude-sonnet \
  --context "You are a financial writing assistant. Draft a follow-up
             email based on this meeting summary: {sanitized_output}" \
  --no-history

# Step 3: Return output to local layer
# Advisor manually substitutes [NAME_1] → actual name before sending
```

What a "PII Wall" Looks Like to a Compliance Auditor

- **Documented policy:** "Client PII is processed only by on-device models. Cloud APIs receive only redacted summaries." (Written in compliance manual)
- **Technical enforcement:** Zone A sessions use Ollama exclusively; Zone B sessions are spawned fresh with sanitized inputs
- **Audit trail:** Session logs show local Ollama calls for PII-handling steps, cloud API calls for drafting steps
- **Vendor agreements:** DPAs executed with Anthropic and OpenAI
- **Periodic review:** Quarterly review of memory embedding contents to confirm no PII leaked through

> 📋 ***For Spectrum:*** *"The 'clean-room' architecture means we can use the most capable AI models in the world for drafting and research without those models ever processing a client's real name or account number — the PII stays on your machine, the intelligence comes from the cloud."*

---

9. Recommended Action Items

Numbered by priority. Items 1–4 are blockers before any client-data workflow goes live.

🔴 Priority 1 — Complete Before Client Deployment

1. Execute Anthropic DPA

- Contact: enterprise@anthropic.com or via Anthropic Console → Settings → Data
- Confirms: no training on API data; establishes formal data processing relationship
- Time estimate: 1–2 business days

2. Execute OpenAI DPA

- URL: https://openai.com/policies/data-processing-addendum
- Consider requesting Zero Data Retention (ZDR) for maximum compliance posture
- Also covers: DALL-E image gen, Whisper transcription, embeddings
- Time estimate: 1–3 business days for enterprise agreements

3. Audit and Purge Memory Embeddings

- Review current embedding index (OpenAI embeddings pipeline)
- Identify any embeddings that may contain client references from Deacon's personal use
- Either purge and rebuild the index, or replace OpenAI embeddings with `nomic-embed-text` running locally via Ollama
- Time estimate: 2–4 hours of technical work

4. Replace OpenAI Embeddings with Local Model

- Install `nomic-embed-text` via Ollama (small, fast, runs on M-series hardware easily)
- Update memory search pipeline to use local embeddings endpoint instead of `api.openai.com/v1/embeddings`
- This eliminates the highest-risk secondary data stream
- Time estimate: 4–8 hours of development work

🟡 Priority 2 — Complete Within 30 Days

5. Execute Twilio DPA

- URL: https://www.twilio.com/en-us/legal/data-protection-addendum
- Needed before using the voice number for any client-related calls
- Consider whether call recording is needed for RIA supervisory compliance

6. Implement Clean-Room Session Tagging

- Add a flag or session type in OpenClaw config to mark sessions as "client-data" vs. "research"
- "Client-data" sessions: use Ollama models only
- "Research" sessions: cloud models allowed, no client identifiers
- Time estimate: 2–4 hours of config work

7. Contact Brave Search API Support on Retention Policy

- Request written confirmation of query log retention period
- Ask whether a DPA/data processing agreement is available for API customers
- Needed for vendor documentation file

8. Monitor xAI for DPA Availability

- Set a quarterly review reminder to check xAI's enterprise documentation
- Until DPA is available: restrict xAI API to non-client-adjacent tasks only
- Contact: enterprise@x.ai for current status

🟢 Priority 3 — Ongoing / Maintenance

9. Draft Internal AI Use Policy for Spectrum

- One-page document covering: which tools are approved, what data can go to cloud models, what must stay local, who reviews AI-drafted communications
- This anchors the compliance manual entry for AI tooling
- Offering to draft this is a strong pitch move

10. Quarterly Data Flow Review - Review this audit document quarterly for policy changes (all vendor policies can change) - Check that DPAs are still current - Verify no new API integrations were added without going through the review process

11. Check Spectrum's ADV Part 1 on SEC EDGAR - Confirm whether SEC-registered or state-registered (shapes compliance requirements) - Review existing technology vendor disclosures for framing

12. Establish Vendor File for AI Tools - File location: firm compliance folder - Contents: DPA copies, policy screenshots with dates, this audit document - Review cycle: annual (or upon material change)

---

10. Data Flow Map Summary

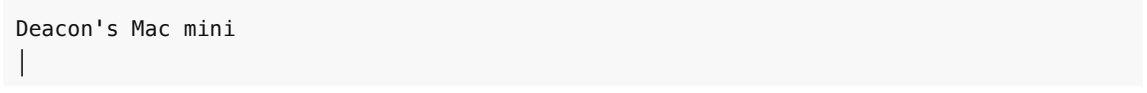

```
├── LOCAL LAYER (zero internet)
│   ├── Ollama: phi4:14b, qwen2.5-coder:14b, gpt-oss:20b, qwen3:8b
│   ├── File system: all client data stored here
│   ├── Local Whisper: transcription on-device
│   └── Memory store: (currently using OpenAI — should be local)
│
├── CLOUD LLM LAYER (data transits, not stored)
│   ├── Anthropic api.anthropic.com — context window per request
│   │   └── DPA available ✅ | No training ✅ | 30-day retention 🟡
│   ├── OpenAI api.openai.com — context + embeddings
│   │   └── DPA available ✅ | No training ✅ | ZDR available ✅
│   └── xAI api.x.ai — context window per request
│       └── No training ✅ | DPA: in progress 🟡
│
├── COMMUNICATIONS LAYER
│   ├── Telegram: bot messages transit Telegram servers 🟡
│   ├── Twilio: call metadata, optional recording
│   │   └── DPA available ✅ | SOC2 ✅ | Recording: opt-in
│   └── ElevenLabs: TTS text (watch what gets passed) 🟡
│
└── SEARCH / LOOKUP LAYER
    ├── Brave Search API: query strings logged to account 🟡
    ├── Google OAuth (NOT YET ACTIVE): email/calendar/Drive 🟡
    └── Google Places: location queries only 🟢
```

Sources: Live config audit of INTEGRATIONS.md and Ollama model list (2026-02-26); Anthropic API Terms of Service and Privacy Policy; OpenAI Enterprise Privacy Commitments and DPA documentation; Twilio Data Processing Addendum documentation; xAI API Terms of Service; Brave Search API documentation; published security hardening review (security-hardening-review_2026-02-16.md); prior spectrum pitch research (spectrum-presentation-script_2026-02-15.md, data-flow-audit-spectrum-2026-02-26.md).

Confidence level: HIGH on local infrastructure facts (directly audited). HIGH on Anthropic/OpenAI/Twilio policies (well-established, publicly documented). MEDIUM on xAI and Brave API policies (newer or less-documented). LOW on Spectrum-specific compliance requirements (not yet gathered).

Prepared: 2026-02-26 | Next review: 2026-05-26